

Eric Yarger, Predictive Analysis

```
In [1]: # Import Libraries

import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
import missingno as msno
from scipy import stats
from scipy.stats import zscore
```

```
In [2]: # Jupyter environment version
!jupyter --version
```

```
jupyter core      : 4.6.3
jupyter-notebook  : 6.0.3
qtconsole         : 4.7.2
ipython           : 7.13.0
ipykernel         : 5.1.4
jupyter client    : 6.1.2
jupyter lab       : 1.2.6
nbconvert         : 5.6.1
ipywidgets        : 7.5.1
nbformat          : 5.0.4
traitlets         : 4.3.3
```

```
In [3]: # Python Environment version
import platform
print(platform.python_version())
```

3.7.7


Data Preparation

Step 1 : Load the Data and initial visualization

```
In [4]: df = pd.read_csv('C:/Users/ericY/Desktop/medical_clean.csv')
```

```
In [5]: msno.matrix(df)
```

```
Out[5]: <matplotlib.axes._subplots.AxesSubplot at 0x1925b98e888>
```




```
In [ ]:
```


Step 2 : Rename columns and create Pairplots


```
In [6]: df.rename(columns={'Item1':'Timely_admis','Item2':'Timely_treat','Item3':'Timely_vis','Item4':'Reliability','It
```

```
In [7]: sns.pairplot(df, x_vars=['TotalCharge'], y_vars=['CaseOrder','Initial_days','VitD_levels','Doc_visits','vitD_su
```


```
Out[7]: <seaborn.axisgrid.PairGrid at 0x1925cef2f88>
```








```
In [8]: sns.pairplot(df, x_vars=['TotalCharge'], y_vars=['Timely_admis', 'Timely_treat', 'Timely_vis', 'Reliability', 'Opti
Out[8]: <seaborn.axisgrid.PairGrid at 0x1925d920748>
```





Step 3 : Address missing data, duplicates, and outliers.


```
In [9]: # Calculate Z-scores, remove Outliers Z > 3
#df.isnull().sum()
```

```
In [10]: df.duplicated().any()
```

```
Out[10]: False
```


```
In [11]: sns.boxplot(df['Income'])
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f076948>
```




```
In [12]: sns.boxplot(df['vitD_supp'])
```

```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f0dc248>
```




```
In [13]: sns.boxplot(df['VitD_levels'])
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f142688>
```




```
In [14]: sns.boxplot(df['Doc_visits'])
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f1abb88>
```




```
In [15]: sns.boxplot(df['Full_meals_eaten'])
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f217648>
```




```
In [16]: sns.boxplot(df['Initial_days'])
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f289c88>
```




```
In [17]: sns.boxplot(df['TotalCharge'])
```

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f2d8808>
```




```
In [18]: sns.boxplot(df['Additional_charges'])
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f344208>
```




```
In [19]: sns.boxplot(df['Age'])
```

```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f3ad5c8>
```




```
In [20]: sns.boxplot(df['Children'])
```

```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f414048>
```



```
In [21]: sns.boxplot(df['Population'])
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f479d08>
```



```
In [ ]:
```

```
In [22]: # Outlier removal method via Z-score, Code reference (Bushmanov, 2019)
```


```
num_data = df.select_dtypes(include=['number'])
cat_data = df.select_dtypes(exclude=['number'])
```

```
In [23]: idx = np.all(stats.zscore(num_data) < 3, axis=1)
```

```
In [24]: df = pd.concat([num_data.loc[idx], cat_data.loc[idx]], axis=1)
```


```
In [25]: sns.boxplot(df['Income'])
```

```
Out[25]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f4fcc48>
```


```
In [26]: sns.boxplot(df['vitD_sup'])
```

```
Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f8aca08>
```




```
In [27]: sns.boxplot(df['VitD_levels'])
```

```
Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f920c88>
```



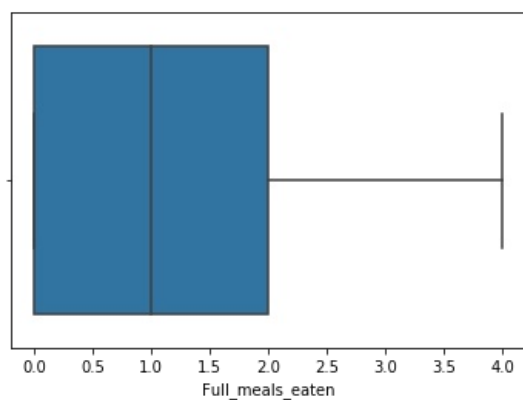
```
In [28]: sns.boxplot(df['Doc_visits'])
```

```
Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f9acdc8>
```




```
In [29]: sns.boxplot(df['Full_meals_eaten'])
```

```
Out[29]: <matplotlib.axes._subplots.AxesSubplot at 0x1925f9f6f48>
```




```
In [30]: sns.boxplot(df['Initial_days'])
```

```
Out[30]: <matplotlib.axes._subplots.AxesSubplot at 0x1925fa71088>
```




```
In [31]: sns.boxplot(df['TotalCharge'])
```

```
Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x1925fae44c8>
```




```
In [32]: sns.boxplot(df['Additional_charges'])
```

```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x1925fb4b488>
```




```
In [33]: sns.boxplot(df['Age'])
```

```
Out[33]: <matplotlib.axes._subplots.AxesSubplot at 0x19260da1708>
```




```
In [34]: sns.boxplot(df['Children'])
```

```
Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x19260e0c408>
```



```
In [35]: sns.boxplot(df['Population'])
```

```
Out[35]: <matplotlib.axes._subplots.AxesSubplot at 0x19260e7b288>
```



```
In [36]: # replace yes/no with 1/0 ReAdmis
df.replace(('Yes', 'No'), (1, 0), inplace=True)
```

```
In [37]: df.duplicated().sum()
```

```
Out[37]: 0
```

Step 4 : Look at correlation between variables

```
In [38]: df.corr()
```


Out[38]:

	CaseOrder	Zip	Lat	Lng	Population	Children	Age	Income	VitD_levels	Doc_visits	...	Str	
	CaseOrder	1.000000	0.010465	-0.012946	-0.012081	0.001489	0.017027	-0.003011	-0.012265	-0.015026	-0.006920	...	0.000
	Zip	0.010465	1.000000	-0.084258	-0.913573	0.012947	0.014307	-0.003327	0.010507	-0.010747	0.000257	...	0.004
	Lat	-0.012946	-0.084258	1.000000	0.001062	-0.187334	0.005874	-0.000132	-0.015414	-0.005158	0.004689	...	-0.001
	Lng	-0.012081	-0.913573	0.001062	1.000000	-0.018263	-0.014141	0.002780	-0.008175	0.000931	0.002417	...	-0.007
	Population	0.001489	0.012947	-0.187334	-0.018263	1.000000	0.007810	-0.018884	0.002162	0.004719	0.016088	...	0.005
	Children	0.017027	0.014307	0.005874	-0.014141	0.007810	1.000000	0.006050	0.003951	0.006542	-0.003467	...	-0.000
	Age	-0.003011	-0.003327	-0.000132	0.002780	-0.018884	0.006050	1.000000	-0.003218	0.008795	0.010819	...	0.011
	Income	-0.012265	0.010507	-0.015414	-0.008175	0.002162	0.003951	-0.003218	1.000000	-0.015684	0.011179	...	0.007
	VitD_levels	-0.015026	-0.010747	-0.005158	0.000931	0.004719	0.006542	0.008795	-0.015684	1.000000	0.010297	...	0.002
	Doc_visits	-0.006920	0.000257	0.004689	0.002417	0.016088	-0.003467	0.010819	0.011179	0.010297	1.000000	...	-0.005
Full_meals_eaten		-0.020805	0.013077	-0.001353	-0.013120	-0.025711	-0.005112	0.008499	-0.012628	0.032606	-0.004586	...	0.002
	vitD_supp	0.026011	0.009348	0.005225	-0.001817	0.004134	-0.010125	0.009336	0.001478	-0.015671	0.002755	...	0.006
	Initial_days	0.831426	0.011103	-0.009938	-0.006659	0.004435	0.022122	0.009943	-0.006543	-0.007267	-0.008363	...	-0.006
	TotalCharge	0.821397	0.010493	-0.012843	-0.005866	0.004758	0.022909	0.010785	-0.008523	-0.004403	-0.005363	...	-0.007
Additional_charges		-0.003178	0.001545	-0.001433	0.003290	-0.011835	0.014076	0.716409	-0.005190	0.006120	0.014611	...	0.033
	Timely_admis	-0.016607	-0.008630	0.008075	0.011933	0.004194	0.004097	0.005614	-0.004194	0.010499	0.003984	...	0.002
	Timely_treat	-0.005508	-0.002475	0.009184	-0.002521	0.016837	0.006169	0.004382	-0.012371	0.003697	0.004377	...	-0.008
	Timely_vis	-0.006320	-0.010277	0.010924	0.002614	-0.004754	-0.002485	0.006990	-0.007394	-0.011930	-0.003794	...	0.002
	Reliability	-0.016204	0.001231	-0.011577	0.000283	-0.008892	-0.001091	0.003407	-0.003532	-0.016650	-0.006303	...	-0.013
	Options	-0.004709	0.006290	0.000179	-0.002771	0.013720	0.003409	-0.013980	-0.005088	0.007878	-0.011124	...	0.004
	Hours	-0.006087	-0.001406	0.009542	-0.004637	0.007970	-0.002796	0.003434	0.003083	0.004610	0.009226	...	-0.000
	Courteous	0.005102	-0.004203	0.009071	0.002070	0.010529	0.015894	0.009339	0.008516	-0.007461	0.005322	...	-0.001
	Listen	-0.012319	-0.010159	0.004348	0.003871	-0.005522	-0.011509	0.002873	0.020238	-0.024347	0.006145	...	0.001
	ReAdmis	0.661462	0.009519	-0.012324	-0.004241	0.007563	0.023890	0.011880	-0.008669	0.002858	-0.002226	...	-0.004
	Soft_drink	-0.014014	-0.000717	0.010164	-0.000521	0.001137	0.017635	0.000994	0.001186	0.009452	0.013203	...	0.004
	HighBlood	0.005010	0.004015	-0.003431	0.002898	-0.000321	0.006837	0.008265	-0.003876	0.004668	0.012391	...	0.005
	Stroke	0.000329	0.004496	-0.001691	-0.007993	0.005435	-0.000677	0.011657	0.007790	0.002774	-0.005126	...	1.000
	Overweight	-0.015372	0.008915	-0.003979	-0.012112	0.007289	-0.017994	-0.009858	-0.018495	0.010462	0.010228	...	-0.002
	Arthritis	0.008205	0.010427	0.011206	-0.014313	-0.006209	0.008645	0.006977	-0.006684	0.006214	0.000778	...	-0.017
	Diabetes	-0.006454	-0.009222	0.022006	0.004272	-0.009324	0.012099	0.007367	-0.010296	-0.024080	0.012104	...	0.011
	Hyperlipidemia	-0.006078	0.004892	-0.003521	-0.004862	-0.003931	-0.001939	0.004963	-0.003735	-0.008225	-0.026669	...	-0.009
	BackPain	0.012056	-0.010125	-0.008733	-0.000215	0.011799	-0.013339	0.026304	0.003499	-0.008061	0.012149	...	0.007
	Anxiety	0.014668	0.007898	0.015698	-0.012313	-0.002665	0.002848	0.001124	0.003670	0.006645	0.002448	...	-0.010
	Allergic_rhinitis	-0.001750	-0.017582	0.013252	0.012777	0.008393	-0.015710	0.013334	-0.004113	-0.005865	0.001936	...	-0.003
Reflux_esophagitis		0.010788	-0.002543	-0.013755	0.000970	0.004870	-0.007773	-0.018236	0.014447	-0.008863	-0.006894	...	0.004
	Asthma	-0.015245	0.010517	0.000171	-0.006567	0.002971	0.002161	0.011092	0.013542	0.000382	-0.016877	...	0.000

36 rows × 36 columns



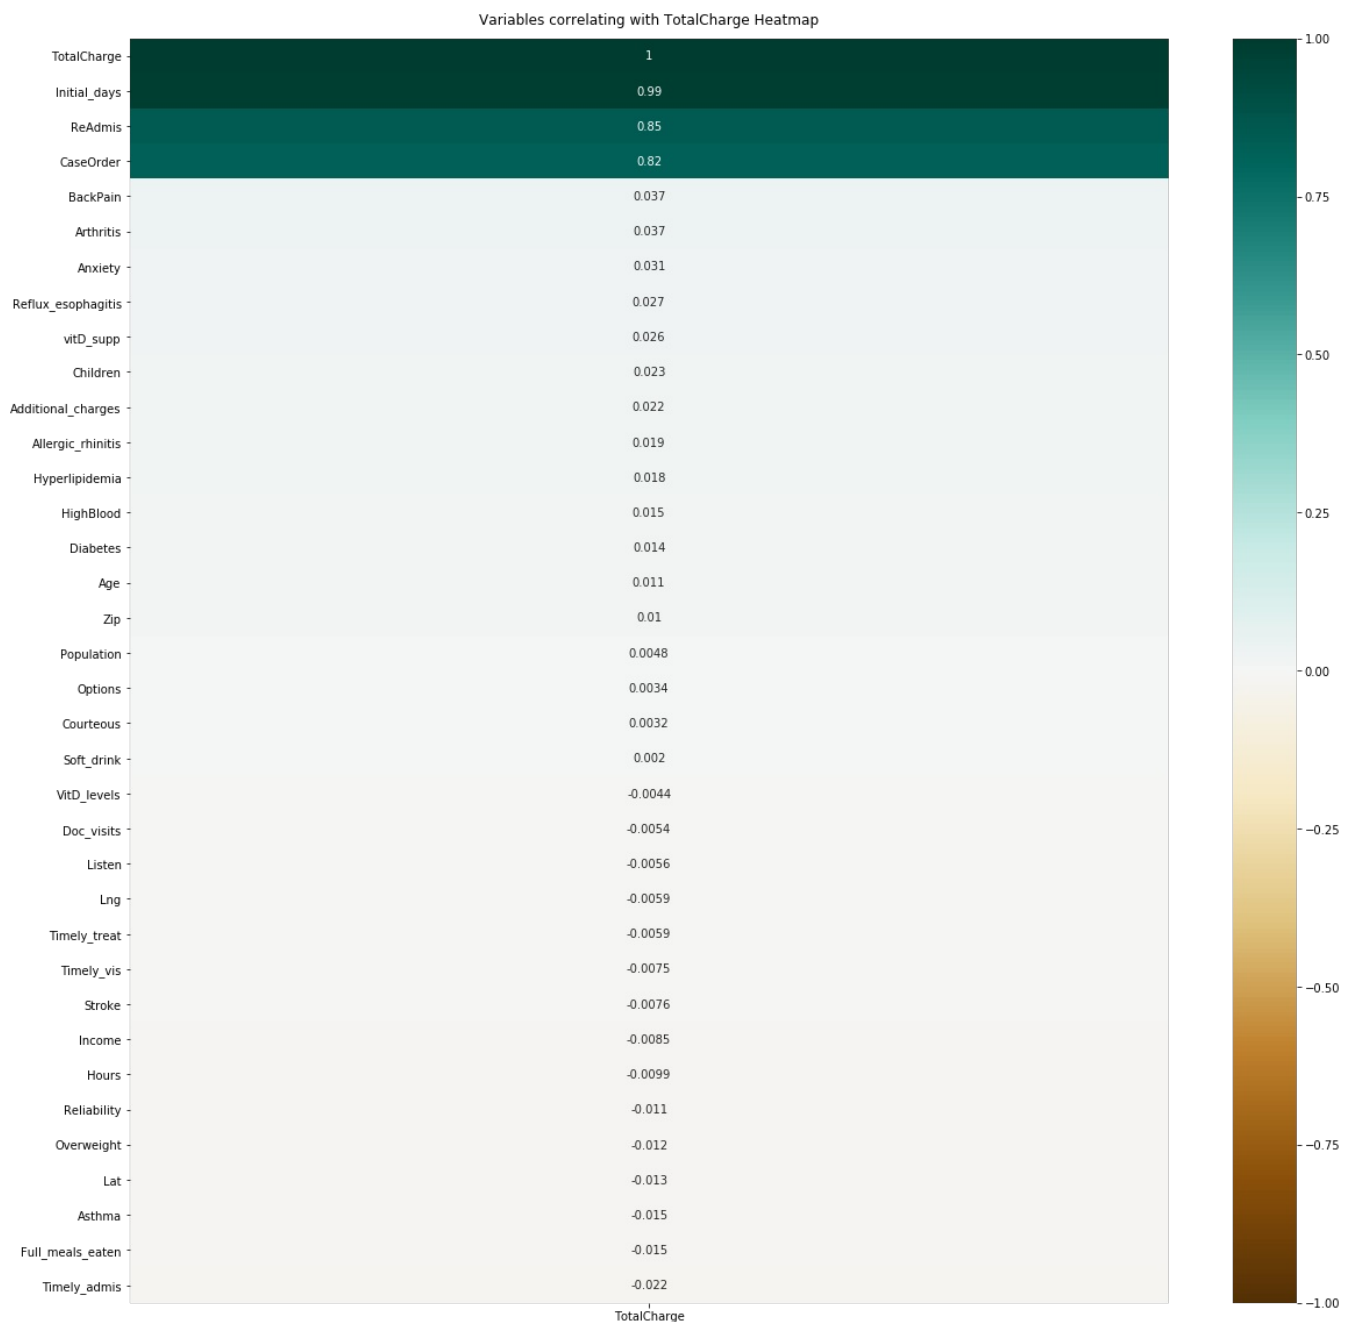
```
In [39]: fig_dims = (20, 20)
fig, ax = plt.subplots(figsize=fig_dims)
sns.heatmap(df.corr(), ax=ax)
plt.show()
```



```
In [40]: # Heatmap code reference (Seaborn.heatmap, n.d.)
```

```
plt.figure(figsize=(20,20))
heatmap = sns.heatmap(df.corr()[['TotalCharge']].sort_values(by='TotalCharge', ascending=False), vmin=-1, vmax=
heatmap.set_title('Variables correlating with TotalCharge Heatmap',pad=12)
```

```
Out[40]: Text(0.5, 1, 'Variables correlating with TotalCharge Heatmap')
```



Step 5 : Create dummy variables (ensure n = k number of variables) & rename any necessary features

```
In [41]: #Get dummies code reference (Pandas.get_dummies, N.d.)
df = pd.get_dummies(df, columns=['ReAdmis', 'Children', 'Marital', 'Gender', 'Doc_visits', 'vitD_supp', 'Soft_drink',

In [42]: df.rename(columns={'Services_CT Scan': 'Services_CT_Scan', 'Marital_Never Married': 'Marital_Never_Married', 'Initi
```

Step 6 drop demographic features that won't be used in the analysis

```
In [43]: df.drop(['CaseOrder', 'Customer_id', 'Interaction', 'UID', 'City', 'State', 'County', 'Area', 'TimeZone', 'Job', 'Lng', 'L
```

Step 7 select features for regression statistically - abs > .04 correlation with TotalCharge

```
In [44]: abs(df.corr()["TotalCharge"][abs(df.corr()["TotalCharge"]>0.04).drop('TotalCharge')).index.tolist()
```

```
Out[44]: ['Initial_days',  
'ReAdmis_0',  
'ReAdmis_1',  
'Initial_admin_Elective_Admission',  
'Initial_admin_Emergency_Admission',  
'Initial_admin_Observation_Admission',  
'Complication_risk_High',  
'Complication_risk_Medium']
```

```
In [45]: df = df[['TotalCharge', 'Initial_days',  
'ReAdmis_0',  
'ReAdmis_1',  
'Initial_admin_Elective_Admission',  
'Initial_admin_Emergency_Admission',  
'Initial_admin_Observation_Admission',  
'Complication_risk_High',  
'Complication_risk_Medium']]  
df
```

```
Out[45]:
```

	TotalCharge	Initial_days	ReAdmis_0	ReAdmis_1	Initial_admin_Elective_Admission	Initial_admin_Emergency_Admission	Initial_admin_C
0	3726.702860	10.585770	1	0	0	1	
1	4193.190458	15.129562	1	0	0	1	
2	2434.234222	4.772177	1	0	1	0	
3	2127.830423	1.714879	1	0	1	0	
4	2113.073274	1.254807	1	0	1	0	
...	
9995	6850.942000	51.561220	1	0	0	1	
9996	7741.690000	68.668240	0	1	1	0	
9997	8276.481000	70.154180	0	1	1	0	
9998	7644.483000	63.356900	0	1	0	1	
9999	7887.553000	70.850590	0	1	0	0	

9206 rows × 9 columns

Step 8 Min-Max scaling of features

```
In [ ]:
```

```
In [46]: df = (df - df.min()) / (df.max() - df.min())  
df
```

```
Out[46]:
```

	TotalCharge	Initial_days	ReAdmis_0	ReAdmis_1	Initial_admin_Elective_Admission	Initial_admin_Emergency_Admission	Initial_admin_C
0	0.246933	0.135022	1.0	0.0	0.0	1.0	
1	0.311343	0.199037	1.0	0.0	0.0	1.0	
2	0.068475	0.053117	1.0	0.0	1.0	0.0	
3	0.026168	0.010044	1.0	0.0	1.0	0.0	
4	0.024130	0.003562	1.0	0.0	1.0	0.0	
...	
9995	0.678314	0.712308	1.0	0.0	0.0	1.0	
9996	0.801304	0.953321	0.0	1.0	1.0	0.0	
9997	0.875146	0.974256	0.0	1.0	1.0	0.0	
9998	0.787882	0.878492	0.0	1.0	0.0	1.0	
9999	0.821444	0.984067	0.0	1.0	0.0	0.0	

9206 rows × 9 columns

```
In [47]: #Read out prepared data set for submission.
```



```
In [48]: df.to_excel('C:/Users/ericy/Desktop/d209.1_prepared.xlsx')
```

```
In [ ]:
```

Step 9 Summary stats for selected features and TotalCharge

```
In [49]: df.isnull().sum()
```

```
Out[49]: TotalCharge          0
Initial_days          0
ReAdmis_0             0
ReAdmis_1             0
Initial_admin_Elective_Admission  0
Initial_admin_Emergency_Admission  0
Initial_admin_Observation_Admission  0
Complication_risk_High  0
Complication_risk_Medium  0
dtype: int64
```

```
In [50]: df.describe()
```

```
Out[50]:
```

	TotalCharge	Initial_days	ReAdmis_0	ReAdmis_1	Initial_admin_Elective_Admission	Initial_admin_Emergency_Admission	Initial_admin_Observation_Admission
count	9206.000000	9206.000000	9206.000000	9206.000000	9206.000000	9206.000000	9206.000000
mean	0.465055	0.470530	0.633065	0.366935	0.250489	0.505323	0.244189
std	0.301177	0.370886	0.481995	0.481995	0.433318	0.499999	0.335108
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.171211	0.096921	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.436588	0.420396	1.000000	0.000000	0.000000	1.000000	0.000000
75%	0.762208	0.847510	1.000000	1.000000	1.000000	1.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

```
In [51]: df.corr()
```

```
Out[51]:
```

	TotalCharge	Initial_days	ReAdmis_0	ReAdmis_1	Initial_admin_Elective_Admission	Initial_admin_Emergency_Admission	Initial_admin_Observation_Admission
TotalCharge	1.000000	0.987666	-0.845034	0.845034	-0.055349	0.107284	-0.069032
Initial_days	0.987666	1.000000	-0.852064	0.852064	0.011339	-0.010895	0.001243
ReAdmis_0	-0.845034	-0.852064	1.000000	-1.000000	0.010482	-0.019393	0.011997
ReAdmis_1	0.845034	0.852064	-1.000000	1.000000	-0.010482	0.019393	-0.011997
Initial_admin_Elective_Admission	-0.055349	0.011339	0.010482	-0.010482	1.000000	-0.584290	0.022966
Initial_admin_Emergency_Admission	0.107284	-0.010895	-0.019393	0.019393	-0.584290	1.000000	-0.328595
Initial_admin_Observation_Admission	-0.069032	0.001243	0.011997	-0.011997	0.022966	-0.328595	1.000000
Complication_risk_High	0.081082	-0.008235	0.004294	-0.004294	0.022966	-0.015254	0.000000
Complication_risk_Medium	-0.064676	-0.006597	-0.006905	0.006905	-0.015254	0.000000	0.000000

```
In [52]: df.mean()
```

```
Out[52]: TotalCharge          0.465055
Initial_days          0.470530
ReAdmis_0             0.633065
ReAdmis_1             0.366935
Initial_admin_Elective_Admission  0.250489
Initial_admin_Emergency_Admission  0.505323
Initial_admin_Observation_Admission  0.244189
Complication_risk_High  0.335108
Complication_risk_Medium  0.452857
dtype: float64
```

```
In [53]: df.median()
```

```
Out[53]: TotalCharge          0.436588
Initial_days          0.420396
ReAdmis_0             1.000000
ReAdmis_1             0.000000
Initial_admin_Elective_Admission  0.000000
Initial_admin_Emergency_Admission  1.000000
Initial_admin_Observation_Admission  0.000000
Complication_risk_High  0.000000
Complication_risk_Medium  0.000000
dtype: float64
```

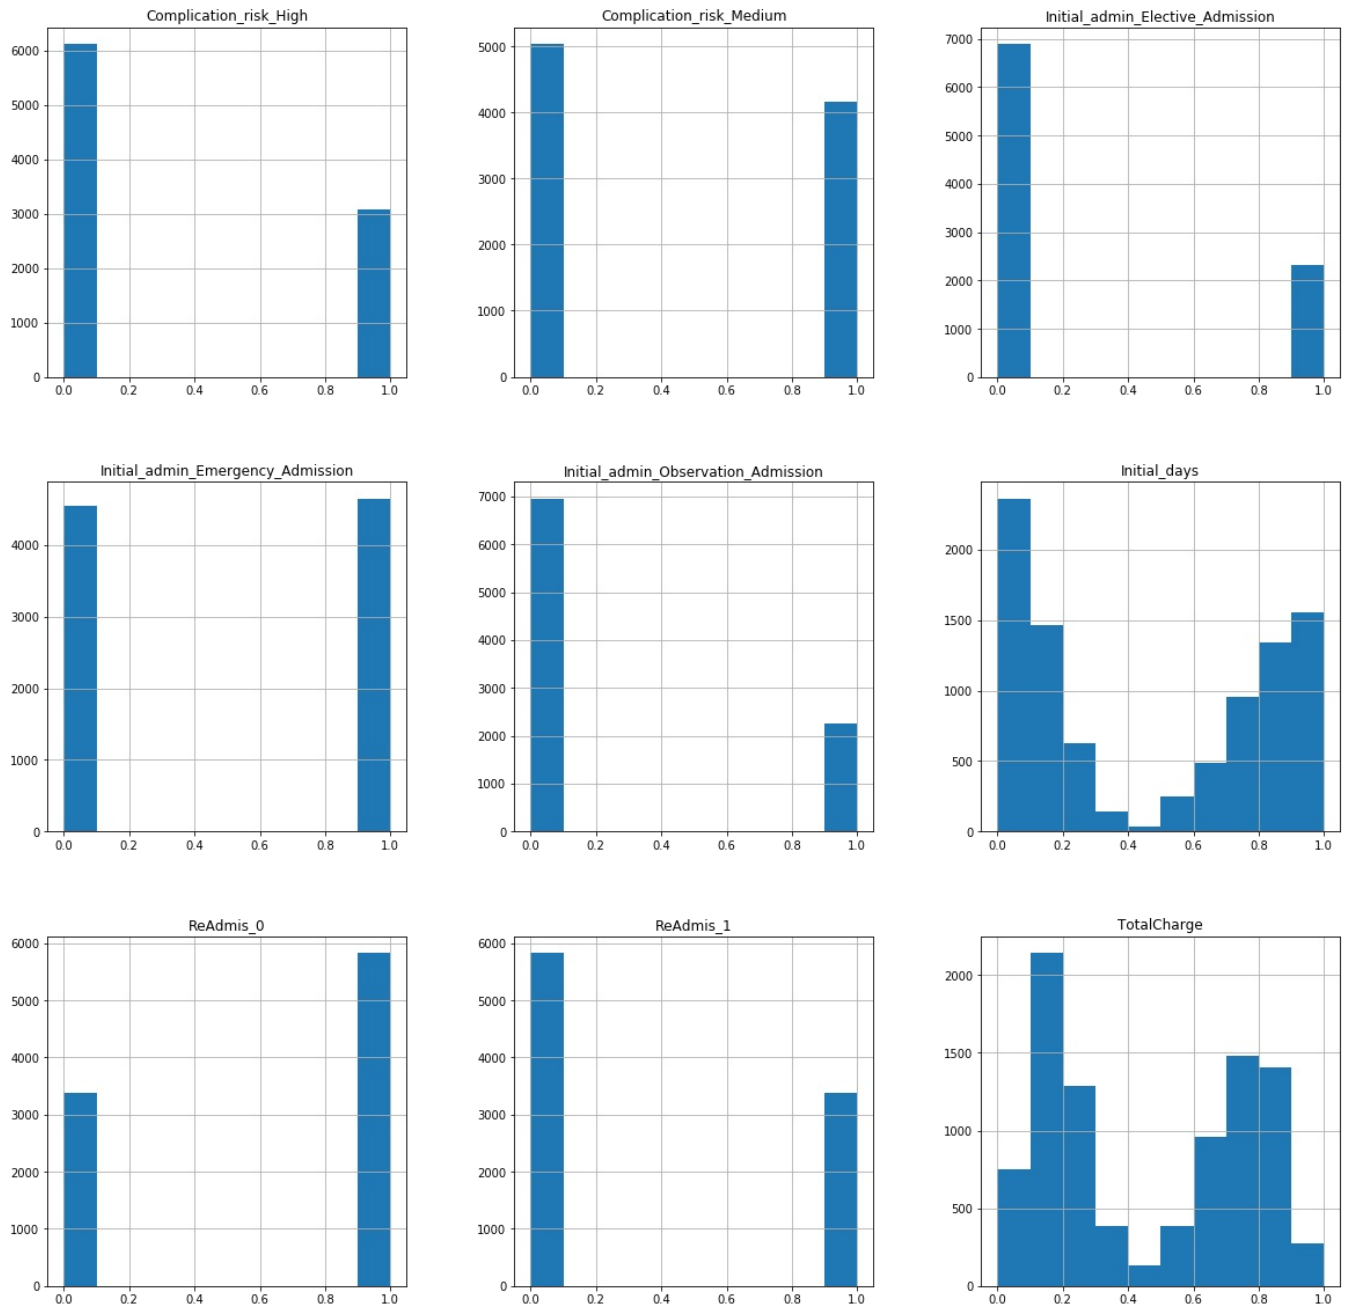
```
In [54]: df.mode()
```

```
Out[54]:
```

	TotalCharge	Initial_days	ReAdmis_0	ReAdmis_1	Initial_admin_Elective_Admission	Initial_admin_Emergency_Admission	Initial_admin_Obse
0	0.775589	0.935755	1.0	0.0	0.0		1.0
1	0.832094	0.976668	NaN	NaN	NaN		NaN

```
In [55]: df.hist(figsize=(20,20))
```

```
Out[55]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x000001925FC95E88>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000019261986B48>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000019261B6EB48>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x00000192620A2C88>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000192623ABD88>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000192623B4EC8>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000019261A87A88>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000192622C3FC8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000192622B7B48>]],
dtype=object)
```



```
In [ ]:
```

Step 10 .astype() categorical features to int8

```
In [56]: df.astype({'ReAdmis_0': 'int8', 'ReAdmis_1': 'int8', 'Initial_admin_Elective_Admission': 'int8', 'Initial_admin_Emergency_Admission': 'int8', 'Initial_admin_Obse': 'int8'})
```

```
Out[56]: TotalCharge          float64
Initial_days          float64
ReAdmis_0             int8
ReAdmis_1             int8
Initial_admin_Elective_Admission  int8
Initial_admin_Emergency_Admission  int8
Initial_admin_Observation_Admission  int8
Complication_risk_High  int8
Complication_risk_Medium  int8
dtype: object
```

```
In [ ]:
```

Step 11 Assign prepared features to y = TotalCharge, X = Prepared independent features.

```
In [57]: y=df['TotalCharge']
X = df.drop(columns = 'TotalCharge')
```

Step 12 import necessary libraries for Decision Tree regression analysis and cross validation
split data into test and train sets.

```
In [ ]:
```

```
In [58]: from sklearn import tree
from sklearn.model_selection import cross_val_score, train_test_split
from sklearn import metrics
from sklearn.metrics import mean_squared_error as MSE
```

Section D Data Analysis

```
In [59]: #Instantiate the model
# Code reference (Instantiate the model, n.d.)
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = .7, test_size = 0.30, random_state = 1)
```

```
In [60]: X_train.to_excel('C:/Users/ericy/Desktop/d209.1.X_train.xlsx')
```

```
In [61]: X_test.to_excel('C:/Users/ericy/Desktop/d209.1.X_test.xlsx')
```

```
In [62]: y_train.to_excel('C:/Users/ericy/Desktop/d209.1.y_train.xlsx')
```

```
In [63]: y_test.to_excel('C:/Users/ericy/Desktop/d209.1.y_test.xlsx')
```

```
In [64]: #Instantiate Decision Tree

regr_1 = tree.DecisionTreeRegressor(max_depth=3, min_samples_leaf=.1, random_state=1)
regr_1.fit(X_train,y_train)
```

```
Out[64]: DecisionTreeRegressor(max_depth=3, min_samples_leaf=0.1, random_state=1)
```

```
In [65]: # Predict

y_pred = regr_1.predict(X_test)
```

Model Accuracy

```
In [66]: # Accuracy of the model
accuracy = regr_1.score(X_test, y_test)
print(accuracy)

0.9618135955188434
```

```
In [67]: #Read predictions and actual y values to variable

ap = pd.DataFrame(data={'Predicted': y_pred, 'Actual': y_test}).head(3000)
ap
```

	Predicted	Actual
5078	0.771395	0.767078
8973	0.647140	0.574817
9724	0.853578	0.774836
4849	0.108177	0.048150
5896	0.771395	0.743521
...
8514	0.647140	0.536069
1704	0.175889	0.134419
7034	0.647140	0.686924
3942	0.108177	0.039914
8391	0.647140	0.723192

2762 rows × 2 columns

```
In [68]: # Calculate Mean Square Error

mse_dt = MSE(y_test, y_pred)
print('MSE: {:.4f}'.format(mse_dt))

MSE: 0.0034
```

```
In [69]: # Compute the array containing the 10-folds CV MSEs
# Code Reference (Evaluate the 10-fold CV error, n.d.)

MSE_CV_scores = - cross_val_score(regr_1, X_train, y_train, cv=10,
                                   scoring='neg_mean_squared_error',
                                   n_jobs=-1)
```

```
In [70]: # Compute Root Mean Square Error from cross val scores

RMSE_CV = (MSE_CV_scores.mean())**(1/2)

# Print RMSE CV
print('CV RMSE: {:.4f}'.format(RMSE_CV))

CV RMSE: 0.0578
```

```
In [71]: # Compute Root Mean Square Error from training set & predictions
# code reference (Evaluate the training error, n.d.)

# Fit dt to the training set
regr_1.fit(X_train, y_train)

# Predict the labels of the training set
y_pred_train = regr_1.predict(X_train)

# Evaluate the training set RMSE of regr_1
RMSE_train = (MSE(y_train, y_pred_train))**(1/2)

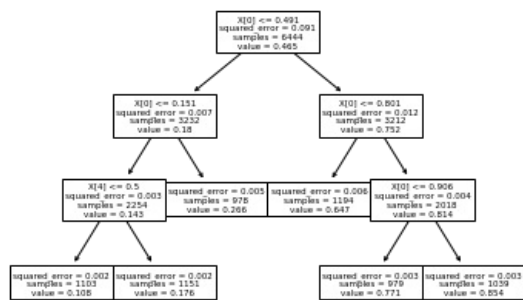
# Print RMSE_train
print('Train RMSE: {:.4f}'.format(RMSE_train))

Train RMSE: 0.0583
```

```
In [ ]:
```


```
In [72]: # Graph decision tree
# Code reference (Galarnyk, 2022)
tree.plot_tree(regr_1)
```

```
Out[72]: [Text(0.5, 0.875, 'X[0] <= 0.491\nsquared_error = 0.091\nsamples = 6444\nvalue = 0.465'),
Text(0.3, 0.625, 'X[0] <= 0.151\nsquared_error = 0.007\nsamples = 3232\nvalue = 0.18'),
Text(0.2, 0.375, 'X[4] <= 0.5\nsquared_error = 0.003\nsamples = 2254\nvalue = 0.143'),
Text(0.1, 0.125, 'squared_error = 0.002\nsamples = 1103\nvalue = 0.108'),
Text(0.3, 0.125, 'squared_error = 0.002\nsamples = 1151\nvalue = 0.176'),
Text(0.4, 0.375, 'squared_error = 0.005\nsamples = 978\nvalue = 0.266'),
Text(0.7, 0.625, 'X[0] <= 0.801\nsquared_error = 0.012\nsamples = 3212\nvalue = 0.752'),
Text(0.6, 0.375, 'squared_error = 0.006\nsamples = 1194\nvalue = 0.647'),
Text(0.8, 0.375, 'X[0] <= 0.906\nsquared_error = 0.004\nsamples = 2018\nvalue = 0.814'),
Text(0.7, 0.125, 'squared_error = 0.003\nsamples = 979\nvalue = 0.771'),
Text(0.9, 0.125, 'squared_error = 0.003\nsamples = 1039\nvalue = 0.854')]
```




In [73]: `plt.plot(ap['Predicted'])`

Out[73]: `<matplotlib.lines.Line2D at 0x1926234eac8>`




In [74]: `plt.plot(ap['Actual'])`

Out[74]: `<matplotlib.lines.Line2D at 0x19261fe8ac8>`




In [75]: `sns.lineplot(data=ap)`

Out[75]: `<matplotlib.axes._subplots.AxesSubplot at 0x19261f944c8>`




In [76]: `sns.lineplot(data=ap['Predicted'])`

Out[76]: `<matplotlib.axes._subplots.AxesSubplot at 0x19261fca8c8>`




```
In [77]: sns.lineplot(data=ap['Actual'])
```

```
Out[77]: <matplotlib.axes._subplots.AxesSubplot at 0x192620703c8>
```



```
In [78]: sns.scatterplot(data=ap['Predicted'])
```

```
Out[78]: <matplotlib.axes._subplots.AxesSubplot at 0x19262151ac8>
```




```
In [79]: # Goodness of fit test, predicted vs actual
```

```
fig, ax = plt.subplots()

ax.scatter(y_test, y_pred, edgecolors=(1, 1, 1))
ax.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'k--', lw=5)

ax.set_xlabel('Actual')
ax.set_ylabel('Predicted')
ax.set_title("Goodness of fit")
plt.show()
```



```
In [ ]:
```

